

Residual Transport Maps for Calibrated Scientific Emulators

Elena V. Petrov^{a,*}, Minh Park^{a,b}

^a*Department of Applied Mathematics, Example University, Boston, 02115, United States*

^b*Center for Scientific Computing, Example University, Boston, 02115, United States*

Abstract

We describe a compiling `elsarticle` manuscript for journal submission. A frozen scientific emulator is calibrated with a monotone residual transport map estimated on a small holdout set. The procedure improves interval coverage on three partial-differential-equation benchmarks without retraining the surrogate. The source uses the standard Elsevier front matter, keywords, and numbered bibliography so it typesets with a conventional TeX Live installation that provides `elsarticle.cls`.

Keywords: scientific machine learning, uncertainty quantification, `elsarticle`, emulator calibration

1. Introduction

Learned surrogates are now common in computational science, but their residual law is rarely calibrated when the operating envelope changes. Journal workflows still expect a familiar `elsarticle` skeleton: title, authors, affiliations, abstract, keywords, numbered sections, and a compact bibliography.

Let f_θ denote a pretrained emulator. We treat calibration as estimating a monotone map on a scalar residual score, then inverting that map to form prediction intervals.

*Corresponding author

Email addresses: `epetrov@example.edu` (Elena V. Petrov), `mpark@example.edu` (Minho Park)

Table 1: Holdout coverage and mean interval width for a 90% target.

Problem	Uncalibrated coverage	Calibrated coverage
Viscous Burgers equation	0.73	0.90
Darcy flow	0.66	0.88
Reaction–diffusion	0.79	0.91

2. Calibration

For a labeled holdout set $\{(x_i, y_i)\}_{i=1}^n$ define $s_i = \|y_i - f_\theta(x_i)\|_2$. Write $\hat{F}$ for the empirical distribution function of $\{s_i\}$. At a query x the symmetric interval of level $1 - \alpha$ is

$$f_\theta(x) \pm \hat{F}^{-1}(1 - \alpha). \quad (1)$$

Exchangeability of calibration and test scores yields the usual conformal $1/(n + 1)$ coverage correction.

3. Related work

Elsevier journals in computational science typically expect a short account of prior art before the method. Physics-informed networks and operator learning supply the frozen emulator [1]. Conformal prediction supplies the coverage argument [2]. We treat residual scores as a one-dimensional interface so that neither ingredient has to be rewritten.

4. Numerical evidence

Table 1 summarizes holdout coverage before and after calibration. Width increases modestly while coverage approaches the 90% target. A second split with $n = 400$ changes coverage by less than 0.03, which suggests that the empirical map is already stable at the smaller holdout size used in the main table.

5. Conclusion

The manuscript is intentionally short: it demonstrates a complete `elsarticle` document that compiles, and it records a practical calibration baseline for scientific emulators.

References

- [1] M. Raissi, P. Perdikaris, G. E. Karniadakis, Physics-informed neural networks, *J. Comput. Phys.* 378 (2019) 686–707.
- [2] V. Vovk, A. Gammerman, G. Shafer, *Algorithmic Learning in a Random World*, Springer, 2005.